

Abhiram Sharma

LLM Infrastructure | Agentic AI Platforms | Machine Learning Engineering
sharmaabhiram1809@gmail.com • +91 79760 90859
linkedin.com/in/sharmaabhiram • github.com/abhiram1809 • abhiram1809.github.io

Experience

Panasonic Avionics

Machine Learning Engineer II (promoted from MLE I, Jun 2026)

Hybrid

Jul 2025 – Present

- Architected an end-to-end LLM routing service for coding agents (Claude Code, Cursor, Kilo Code), improving cost utilization 4x through request classification and model-tier routing.
- Founded and scaled the organization's LLM inference platform serving open-source and closed-source models behind a unified, OpenAI-compatible endpoint that works across LLM SDKs, increasing engineering productivity 60%.
- Built an agentic AI platform that standardized development and deployment of production-grade AI agents, reducing agent creation time from days to under 10 minutes.
- Architected a Kubernetes controller platform (Operators, CRDs) for organization-wide LLM deployments, with automated lifecycle management, budget controls, and centralized observability, cutting debugging and incident resolution time 90%.
- Mentored an engineering intern in AWS, AI-driven engineering practices, and independent production troubleshooting.

Calsoft, Inc.

Data Scientist – Consultant for Panasonic Avionics

Pune, Maharashtra

Jan 2025 – Jul 2025

- Built LLM input/output assessment guardrails using open-source and closed-source technologies, achieving an 88% safer AI system.
- Engineered a WAF-fortified accelerated inference pipeline on vLLM, reaching 80 tokens/second on dense models for proprietary coding and application workflows.

Softsensor.AI

Associate Data Scientist

Jaipur, Rajasthan

Sep 2023 – Jan 2025

- **Cyber Security Chatbot:** Built core Text2SQL and unstructured MITRE document query workflows for precise log and attack-pattern insights.
- **Claims Automation:** Led full-cycle unstructured-to-structured extraction for letters, emails, and SRRs with interlinked document relationships, cutting extraction cost up to 85%.
- **Document Translation:** Built end-to-end translation software preserving tables and diagrams across document formats.

Blackcoffer Consulting

Associate Data Scientist

Remote

Mar 2022 – Feb 2023

- Developed a computer vision solution for satellite imagery analysis, automating detection and localization of vehicles in high-resolution imagery.

Projects & Open Source

- **vLLM – MiniMax-M3 on NVFP4/Blackwell** (PR in review): Root-caused and fixed four parser bugs that garbled output, misrouted streaming reasoning tokens, and dropped tool calls on RTX PRO 6000 (vllm-project/vllm#47001).
- **LLM Fine-Tuning Advisor (open source):** Agent skill for Claude Code, OpenCode, and Codex that plans LoRA, SFT, and continued-pretraining runs, diagnoses local training failures, and evaluates the trained artifact on the vLLM endpoint that will serve it. github.com/abhiram1809/unsloth-finetuning-vllm-skill
- **Open Source – Merged Contributions:** *LlamaIndex* – added Google Sheets ingestion into pandas DataFrames (#11640) and embeddings for custom knowledge-graph triplets (#10629). *HuggingFace lighteval* – documented the hosted vllm provider prefix needed to reach a locally served vLLM endpoint, which otherwise fails silently despite a correct config (#1060).
- **Quant_AgentTools (open source):** Enabled quantized LLMs to run tool-calling workflows on CPU-only systems via multi-threaded inference for low-cost agent execution. github.com/abhiram1809/Quant_AgentTools

Technical Skills

Languages: Python, SQL, JavaScript, C++, Java, Bash

AI & Machine Learning: Large Language Models (LLMs), Agentic AI, Retrieval-Augmented Generation (RAG), Fine-Tuning (LoRA, QLoRA, PEFT), Quantization, Distributed Training, Model Evaluation, Prompt Engineering, Natural Language Processing (NLP), Computer Vision, Multimodal AI, Knowledge Graphs, Text2SQL, AI Guardrails (NeMo Guardrails), Topic Modeling, ColBERT Reranking

Inference & Serving: vLLM, SGLang, Triton Inference Server, llama.cpp, Ollama, Ray, CUDA, GPU & CPU Inference, Inference Optimization

Infrastructure & MLOps: Kubernetes (Operators, CRDs), Helm, ArgoCD, Docker, CI/CD, GitHub Actions, Prometheus, Grafana, LLM Observability, MLflow, WAF, Git, Linux

Cloud & Data: AWS, Microsoft Azure, PostgreSQL, pgvector, Qdrant, Redis, Kafka, Pandas, NumPy, Data Pipelines, ETL

Frameworks & Tools: PyTorch, HuggingFace Transformers, LlamaIndex, LangChain, LangGraph, CrewAI, LiteLLM, Unsloth, FastAPI, gRPC

Education

Amity University

B.Tech, Computer Science

Jaipur, Rajasthan

Graduated Jan 2024

Certifications: Neural Networks and Deep Learning (DeepLearning.AI) • Machine Learning (Coursera) • Microsoft Azure AI-900 Fundamentals